了解IO设备

核心系统数据库组 余锋http://yufeng.info

@淘宝褚霸

2012-03-17

提纲

- 芯片组
- SATA/SAS
- SSD
- PCIe Flash卡
- RAID卡
- NVRAM卡
- 测量工具

IO芯片组

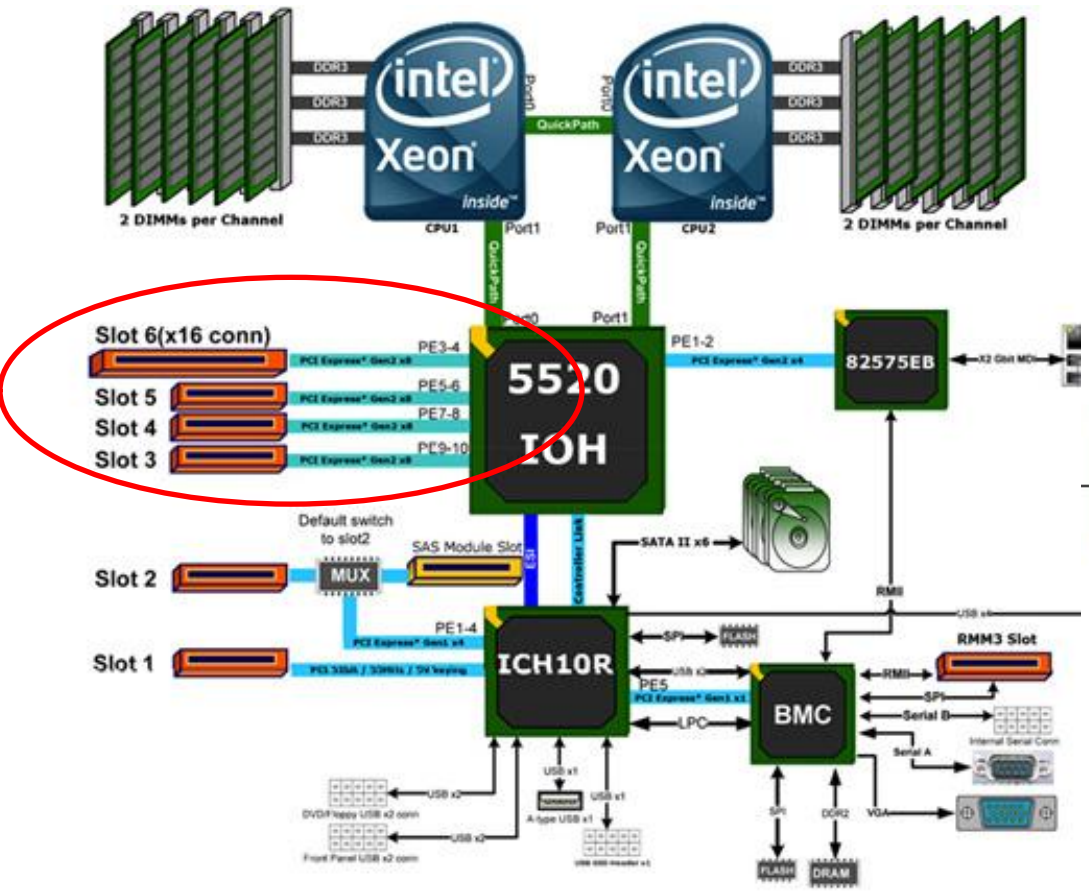

Raid卡和PCIe卡都插在PCIe卡，直接走北桥

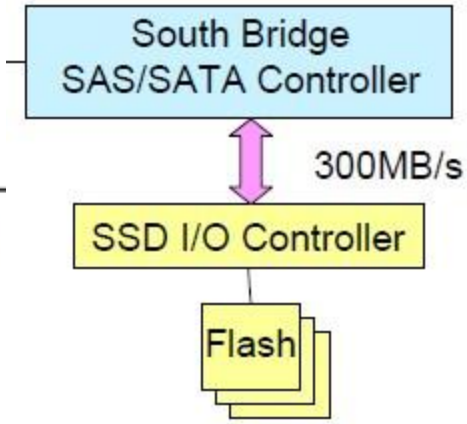

芯片组型号

Processors: 2 x Xeon E5645 2.40GHz 5860MHz FSB

(HT enabled, 12 cores, 24 threads)

Chipset: Intel 5500 IOH-24D B3 (Tylersburg),

82801JIR A0 (ICH10R)

接口速率

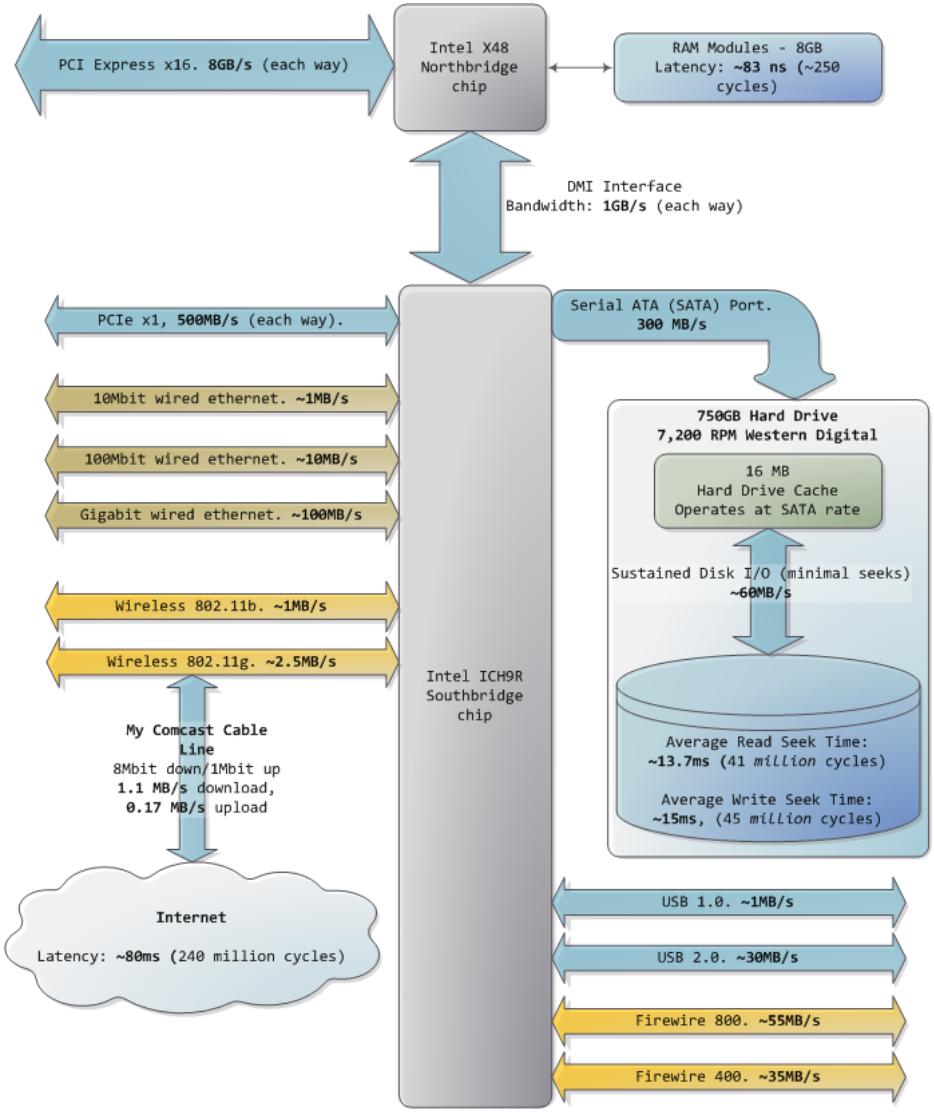

http://duartes.org/gustavo/blog

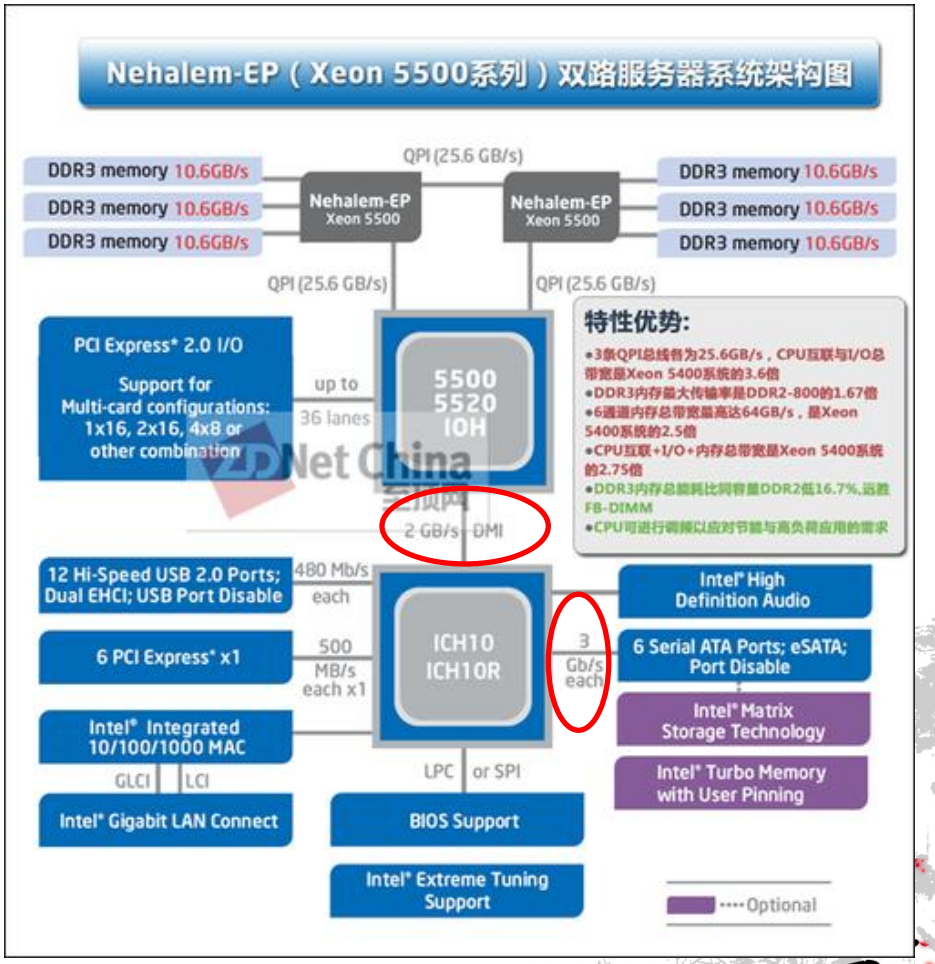

PCIe每个X接口速率:

v1.x: 250 MB/s

v2.x: 500 MB/s

SATA/SAS机械磁盘

SATA II
7200 RPM IOPS: ~90

SAS
15K RPM IOPS: ~180

Disk: sda (scsi0): 100GB JBOD == 1 x HITACHI-HUSSL4010ASS600

SSD

寿命:200T

SATA II
Intel X25-M IOPS: ~8600

为什么

要有RAID或者HBA卡

接SATA磁盘阵列？

解决什么问题？

PCIe Flash卡

PCIe 2.0x8
850 MB/s (4KB)
220,000 IOPS (4KB)

PCIe 2.0x4
ioDrive IOPS: with Flash 140,000
Read IOPS, 135,000 Write IOPS

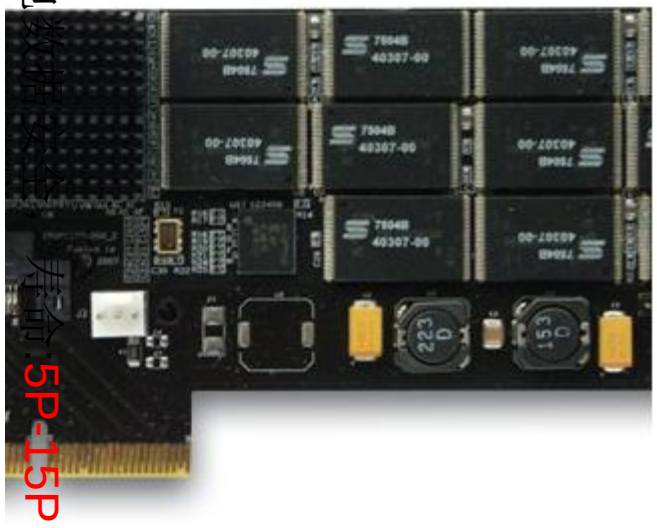

掉电数据保护
接口：5P-15P

Disk-Control: iodrive0: Fusion-io ioDIMM3 320GB

Raid卡

- PCIe 2.0x8
- Support Up to 128 SATA Devices
- Dual Core ROC
- 1 GB cache

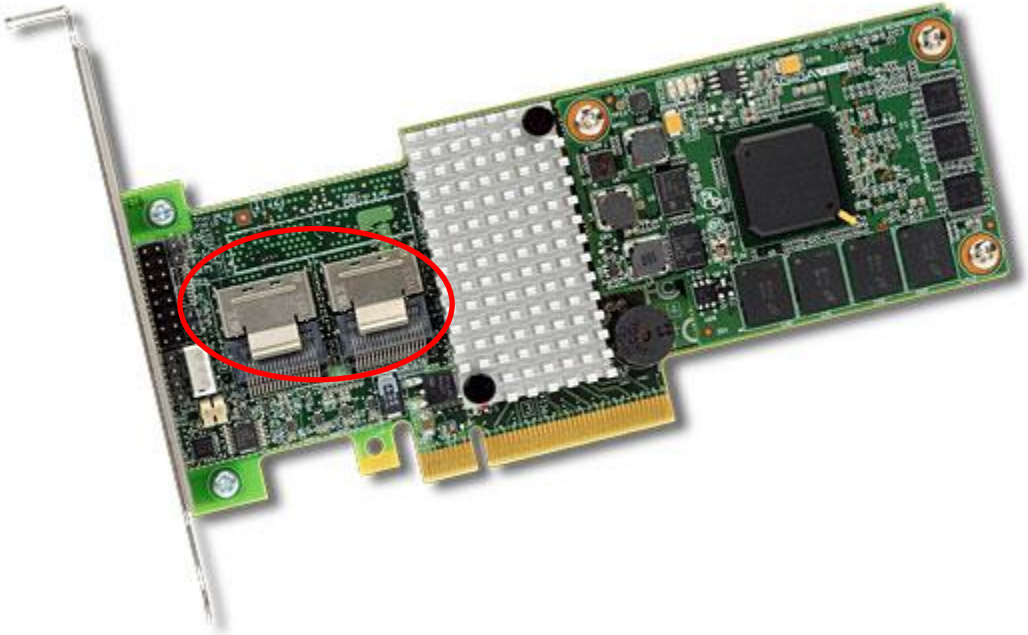

Disk-Control: megaraid_sas0: LSI Logic / Symbios Logic MegaRAID SAS 1078

Raid卡(续)-Cache

- **虚拟卷**
- **预读缓存**
  - NORA (No read ahead)
  - RA (Read ahead)
  - ADRA (Adaptive read ahead)
- **写缓存**
  - WT (Write through),
  - WB (Write back)
- **Disk Cache**
  - 关闭，考虑到数据安全

Raid卡(续)-BBWC

- **Nickel Metal Hydride (NiMH)**
  - 100 full discharge cycles.
  - 48-hour battery life .
  - Typical capacity for the HP Smart Array battery pack reduces by 5 to 10 percent over a 3-year period.
  - Battery recharge takes between 30 minutes and 2 hours
- **模块化设计可替换**

Raid卡 (续) - FBWC

Figure 1. FBWC block diagram

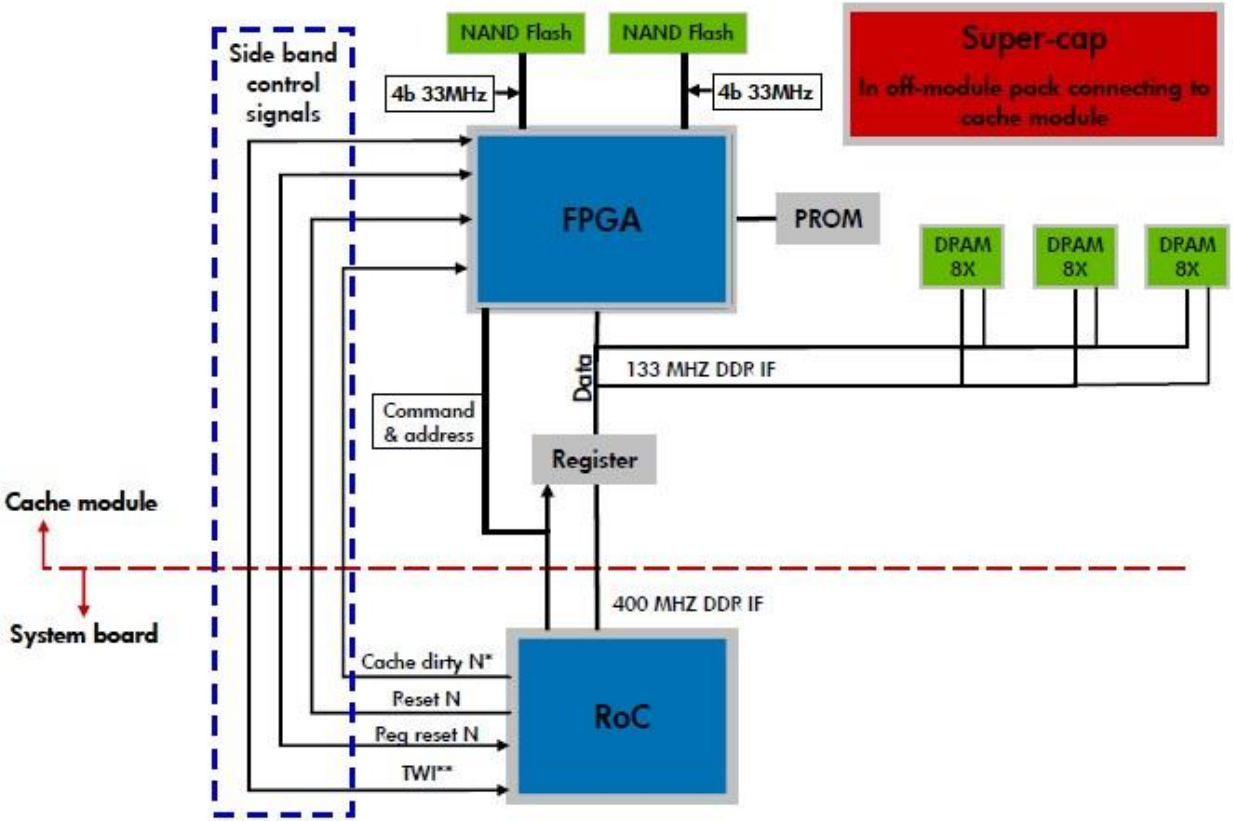

NVRAM卡

寿命:1M hours

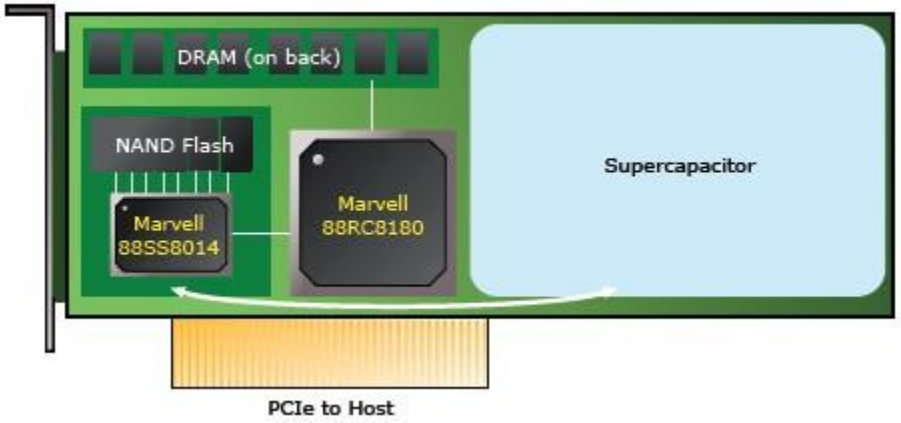

DDR backup to persistent flash on powerfailure
Automatic restore from Flash to DDR when power is restored

PCIe 1.1x4
4K Block Writes: 165,000 IOPS
4K Block Reads: 185,000 IOPS

Disk-Control: mvloki0: Marvell Device 8180

DDR3 Non-Volatile DIMM

static unsigned long ram_start=0xa40000000UL;
static unsigned long ram_size= 0x80000000UL;

PCIe卡的寿命和安全如何保证？

掉电数据安全吗？

hwconfig

hwconfig

```
Disk:           megaraid_sas0-free: JBOD == 10 x 300GB SAS HITACHI-HUS156030VLS600
                                            2 x 300GB SAS HITACHI-HUC106030CSS600
Disk:           sda (scsi0): 599GB (4%) RAID-0 == 2 x 300GB SAS HITACHI-HUS156030VLS600
Disk-Control:   megaraid_sas0: LSI Logic / Symbios Logic MegaRAID SAS 2108 [Liberator],
0.2-0004, FW 2.100.03-1405, BIOS 3.18.00_4.09.05.00_0x0416A000, Cache 1GB, BBU
Disk-Control:   mvloki0: Marvell Device 8180
Chipset:        Intel 5500 IOH-24D B3 (Tylersburg), 82801JIR A0 (ICH10R)
```

hwconfig –x sample.cfg

```
firmware="E516"
handle="69"
interface="SAS"
serial="JXYGHLAN"
size="299999690752"
status="free"
volume="megaraid_sas0-free"
volume_handle="74"
wwn="0x5000cca018c378f1”
model="HITACHI-
HUS156030VLS600"
```

IO子系统架构图

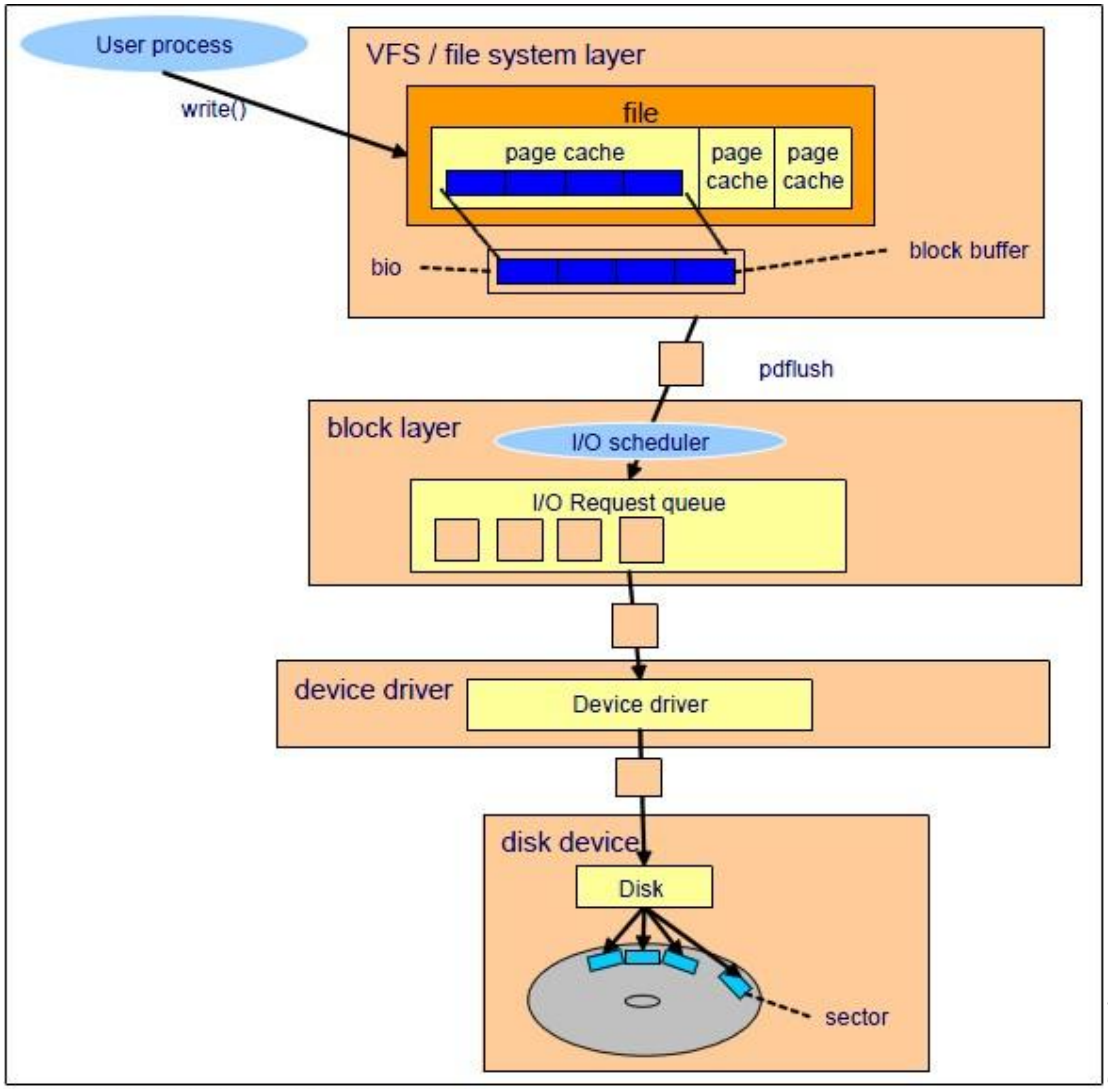

Figure 1-18 I/O subsystem architecture

lsblk

```
# lsblk -i
NAME    MAJ:MIN RM   SIZE RO MOUNTPOINT
sda       8:0    0 557.8G  0
sda1      8:1    0   500M  0 /boot
sda2      8:2    0 146.5G  0 /
sda3      8:3    0     2G  0 [SWAP]
sda4      8:4    0     1K  0
sda5      8:5    0 408.8G  0 /disk0
nvdisk0 252:0    0     8G  0 /u05
```

fio设备写饱和脚本

$ fio a_b_c_d_test

```
[global]
bs=4K
ioengine=libaio
rw=randrw
rwmixwrite=100
time_based
runtime=3600
direct=1
group_reporting
randrepeat=0
norandommap
invalidate=1
iodepth=8
iodepth_batch=4
iodepth_low=4
iodepth_batch_complete=8
numjobs=1

[test_sda]
filename=/dev/sda
[test_sdb]
filename=/dev/sdb
[test_sdc]
filename=/dev/sdc
[test_sdd]
filename=/dev/sdd
```

iostat

Device:	rrqm/s	wrqm/s	r/s	w/s	rsec/s	wsec/s	avgrq-sz	avgqu-sz	await	svctm	%util
sdf	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
sde	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
sda	0.00	0.00	6422.60	0.00	205523.20	0.00	32.00	28.19	4.39	0.16	100.12
sdb	0.00	0.00	5628.60	0.00	180115.20	0.00	32.00	28.21	5.02	0.18	100.12
sdc	0.00	0.00	3316.80	0.00	106137.60	0.00	32.00	28.36	8.55	0.30	100.12
sdd	0.00	0.00	4334.60	0.00	138707.20	0.00	32.00	28.29	6.53	0.23	100.12
memdiska	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00

IO depth对设备性能有什么样的影响？

pidstat

```
root@snb # pidstat -C fio -t -d 1
Linux 2.6.32-131.0.15.el6.x86_64 (snb)  03/17/2012      _x86_64_        (32 CPU)

11:43:35 PM      TGID       TID   kB_rd/s   kB_wr/s kB_ccwr/s  Command
11:43:36 PM     90299         -   2233.96   3200.00      0.00  fio
11:43:36 PM         -     90299   2233.96   3200.00      0.00  |__fio
11:43:36 PM     90300         -   3215.09   3184.91      0.00  fio
11:43:36 PM         -     90300   3215.09   3184.91      0.00  |__fio
11:43:36 PM     90303         -   3230.19   3169.81      0.00  fio
11:43:36 PM         -     90303   3230.19   3169.81      0.00  |__fio
11:43:36 PM     90306         -   3803.77   3200.00      0.00  fio
11:43:36 PM         -     90306   3803.77   3200.00      0.00  |__fio
```

iotop

```
Total DISK READ: 217.69 M/s | Total DISK WRITE: 216.90 M/s
  TID  PRIO  USER     DISK READ  DISK WRITE  SWAPIN      IO>    COMMAND
90300 be/4 root        54.73 M/s   53.43 M/s  0.00 % 91.43 % fio 720-test
90306 be/4 root        53.64 M/s   54.16 M/s  0.00 % 89.22 % fio 720-test
90303 be/4 root        54.99 M/s   54.63 M/s  0.00 % 89.16 % fio 720-test
90299 be/4 root        54.29 M/s   54.66 M/s  0.00 % 89.10 % fio 720-test
```

参考材料

- Fio测试工具使用：
  http://blog.yufeng.info/archives/tag/fio
- hwconfig查看硬件信息:
  http://blog.yufeng.info/archives/2086
- Linux下方便的块设备查看工具lsblk
  http://blog.yufeng.info/archives/1882
- Linux TASK_IO_ACCOUNTING功能以及如何使用:
  http://blog.yufeng.info/archives/2138

谢谢大家！